

Can pass parameters to OS by passing parameters to registers OR passing a block to a register OR pushing/popping onto stack

**Policy:** what will be done? (flexible definition) **Mechanism:** how to do it? (rarely change)

**Process:** active entity loaded into memory for execution and is allocated a processor CPU and other resources

- Memory layout: **text, data, heap, stack**
- States: **new, running, waiting, ready, terminated** (can move from waiting queue to ready queue)
  - Blocked process is put into wait queue
  - Interrupts can remove process and put into ready queue
- Information stored in PCB (maintains process state, program counter, CPU registers, etc)
  - Used when **context switching** from the current process to a new process (process state is pushed on PCB)
- Can be **I/O bound** or **CPU bound**
- **Swapping** can be done to reduce **degree of multiprogramming**
- Parent-Child relationships:
  - OS allocates new resources to child OR child takes a subset of parent resources
  - Parent can concurrently execute with its child OR wait until some or all of its children have terminated
  - Child process can share a copy of parent address space OR load a new program (**exec()**)
  - **Zombie:** process has terminated but parent hasn't called **wait()** **Orphan:** process with no parents
  - When parent terminates, can SOMETIMES result in **cascading termination**

**IPC:** provides **information sharing, computation speedup, and modularity** between **cooperating processes**

- **Shared memory** (read/write to shared area) or **message passing** (send/receive messages from **communication link**)
- **Producer-Consumer** can use **unbounded buffer** or **bounded buffer** (processes maintain **in** and **out** variables)
- **Direction Communication** requires explicit naming, **Indirect Communication** has messages sent to a **mailbox**
  - **Synchronous** (blocks until message is received or available) or **Asynchronous** (receiver can get message or NULL)
  - **Zero, Bounded, or Unbounded Capacity** on the message queue

---

**Threads:** basic unit of CPU utilization that shares data with other threads belonging to the same process

- Shares code section, data section, and OS resources. PCB maintains PC pointers for each thread (**Multithreading**)
- Provides **Responsiveness, Resource Sharing, Economy, and Scalability** (for multicore environment)
- **Concurrency** means thread execution is interleaved **Parallelism** means threads are executed in parallel
  - Types of parallelism: **Data Parallelism** and **Task Parallelism**
- Multithreading models used for when user thread make call to kernel thread
- Thread libraries implemented in user space (uses local function calls) OR in kernel space (uses system calls)
- **Asynchronous Threading** usually involves little data sharing. **Synchronous Threading** usually has a lot of data sharing

**Implicit Threading:** abstracts the complexities of creating and mapping tasks to separate threads for execution

- **Thread Pools:** create finite number of threads that wait for a request (don't have to waste time creating new threads)
- **OpenMP:** programmer identifies regions that can be run in parallel using compiler directives
- **Grand Central Dispatch:** tasks are placed in **serial queue** or **concurrent queue** and are eventually assigned a thread

**Threading Issues:** **fork()/exec()** semantics, **signal handling, thread cancellation, TLS, scheduler activations**

- Receiving signals **Synchronous** (from same process) or **Asynchronously** (from another process)
- Semantics of sending signals to multithreaded programs (all threads or just a few of them)
- **Asynchronous** (immediately) or **Deferred** cancellation (target thread checks itself)
- Scheduler activations use **LWPs** to schedule user thread to run attached to a kthread. Kernel inform apps through **upcalls**

---

Systems consist of manages **safety** (objects across multiple activities) and **liveness** (span across multiple objects)

**Critical-Section Problem:** must satisfy **Mutual Exclusion, Progress, and Bounded Waiting**

Single cores can just disable interrupts during critical section. For multiprocessors, use **preemptive** or **nonpreemptive kernels**

**Peterson's Solution:** manages 2 shared variables: `int turn` and `bool flag[2]` (process is ready to enter CS). Only for 2 processes

```
flag[i] = true;
turn = j;
while (flag[j] && turn == j);    // busy wait

/* critical section */

flag[i] = false;

/* remainder section */
```

**Bakery's Solution:** before CS, each  $P_i$  receives a number.  $P$  with smallest number enters CS. If equal, then  $i < j \implies P_i$  before  $P_j$

**Atomic** instructions either: test word and set value OR swap contents of 2 words

```
bool test_and_set (bool *target) {
    bool rv = *target;
    *target = TRUE;
    return rv;
}

int compare_and_swap(int *value, int expected, int new_value) {
    int temp = *value;

    if (*value == expected) {
        *value = new_value;
    }
    return temp;
}
```

**Mutex:** process calls `acquire()` and `releases()` before entering/leaving critical section. Results in **busy waits**.

**Semaphores:** integer variable accessed through atomic operations `wait()` and `signal()` (**counting** vs **binary**)

- **Note:** moves lock management to critical section (potential issue w/ busy waits), whereas mutex had them before and after CS
- Can force  $P_2$  to execute  $S_2$  only AFTER  $P_1$  executes  $S_1$  (`signal()` called after  $S_1$  and `wait()` called before  $S_2$ )
- Instead of busy waits, can **suspend** process: move it to waiting queue. Can lead to **deadlock**, **starvation**, **priority inversion**

**Monitors:** abstract functions for process synchronization (only 1  $P$  can enter monitor at a time)

- Maintains condition variables:  $P_i$  that invokes `x.wait()` is suspended until `x.signal()` invoked by another  $P_j$
- **Note:** `x.signal()` only affects if another  $P$  is waiting. For semaphore, `signal()` always affected semaphore state
- 2 options of **signal and wait** OR **signal and continue**
- **Conditional wait** can be used to select next  $P$  to resume in monitor where  $P$  with lowest number (highest priority)