

Process: program in execution. Can be **I/O Bound** or **CPU Bound**

- Each process has its own copy of data
- **Cascade termination:** when a parent terminates, its children are forcibly terminated

Activation Record: function params, local vars, return address pushed to stack when a function is called

PCB: contains info about a process. Used in **context switching** when current state saved to PCB and new state is loaded

Interprocess Communication through **shared memory** (shared region of memory) or **message passing** via communication link

- **Direct communication** requires naming recipient/sender. **Indirect communication** has messages sent to **mailboxes**
- Messages might also be buffered

Threads share code section, data section, and OS resources (e.g. open files and signals)

- Challenges include **identifying task**, **load balancing**, **data splitting**, **data dependency**
- **Implicit threading** done through **thread pools**, **fork join**, **OpenMP**, **Grand Central Dispatch**
- **Concurrency** has multiple tasks making progress. **Parallelism** has multiple tasks working simultaneously
- Issue of semantics where should **fork()** copy all threads or just one thread in the context of if **exec()** will be called

Benefits of multithreading: **responsiveness**, **resource sharing**, **economy**, **scalability**

Data parallelism performs same operation on subsets of the data whereas **task parallelism** distributes tasks across multiple cores

User **threads** are mapped to **kernel threads** through **many-to-one**, **one-to-one**, or **many-to-many** models

Asynchronous threading: parent and child threads run concurrently. **Synchronous:** parent waits for children to terminate

Signal: notifies process that an event has occurred, received **synchronously** or **asynchronously**

- For multithreaded programs, can deliver signal to target thread, every thread in process, or thread assigned to receive all signals

Thread cancellation: terminating a thread **asynchronously** (immediately) or **deferred** (let it terminate on its own)

Thread Local Storage (TLS): copy of certain data unique to each thread

Lightweight Process (LWP): schedules user thread to run on an attached kernel thread

- Kernel thread notifies user thread of events through **upcalls**

Cooperating Process affect other processes and can share same logical address space or share data through IPC mechanisms

- May result in **race conditions** where several process manipulate the same data concurrently, creating varying outcomes

Critical section: code that is accessing data shared by other processes. Consists of **entry**, **exit**, and **remainder section**

- Must satisfy **mutual exclusion**, **progress**, **bounded waiting**
- In single-core environment, disable interrupts. For multicore environments, use **preemptive** or **nonpreemptive kernels**
 - **Preemptive** can lead to race conditions. **Nonpreemptive** prevents race conditions from happening

Peterson's Solution: 2 processes share **turn** and **flag** vars: whose turn it is to enter critical section and if the process is ready

Mutex Lock: protects critical sections and prevents race conditions by having processes **acquire()** and **release()** the lock

Busy wait: process that try to enter their critical section are continuously calling **acquire()**, wasting CPU cycles

- **Spin locks** avoid context switching so are preferred for short busy waits

Semaphore: integer variable accessed using **wait()** (decrement) and **signal()** (increment). Either **counting** or **binary semaphore**

- To avoid busy wait, **wait()** can suspend the process if **semaphore <= 0**. It will restart once another process executes **signal()**

Monitor: uses abstract data type to define operations with mutual exclusion and variables that contain the data type state

- Ensures only one process at a time can be active within the monitor

Liveness: properties system must satisfy to ensure progress. Possible issues

- **Deadlock:** multiple processes from a set all waiting for an event from another process in the set
- **Priority Inversion:** higher priority process waiting for a local process to release a resource, but is preempted by another process